

```
# QUESTION 2
```

```
library(ggplot2)
```

```
# Load the dataset from the CSV file
```

```
data <- read.csv("/Users/prabuddhadurge/Downloads/adult_data.csv")
```

```
X <- data$occupation.num
```

```
# Discrete variable
```

```
Y <- data$age
```

```
# Continuous variable
```

```
# (2a) Frequency distribution for discrete variable X (e.g., occupation.num)
```

```
frequency_distribution <- table(X)
```

```
print(frequency_distribution)
```

```
## X
```

```
##      1      2      3      4      5      6      7      8      9     10     11     12     13     14
```

```
## 1350 3721 3212  989 3584 1572 1966 4030  912  644  143    9 3992 4038
```

```
# (2b) Plot the histogram for continuous variable Y
```

```
# Calculate bin width using Freedman-Diaconis rule
```

```
iqr <- IQR(Y)
```

```
# Interquartile range
```

```
bin_width <- 2 * iqr / (length(Y)^(1/3))
```

```
# Freedman-Diaconis rule
```

```
num_bins <- ceiling((max(Y) - min(Y)) / bin_width) # Number of bins based on bin width
```

```
# Plot the histogram
```


```
hist(Y, breaks = num_bins, col = "lightblue",
```

```
      main = "Histogram of Age",
```

```
      xlab = "Age",
```

```
      ylab = "Frequency")
```


## Histogram of Age



```
# (2c) Box-and-Whisker Plot for X and Y
par(mfrow = c(1, 1))


# Boxplot for continuous variable Y
boxplot(Y, main = "Boxplot of Age", col = "lightgreen")
```

## Boxplot of Age




```
#Boxplot for continuous variable X  
boxplot(X, main = "Boxplot of Occupation-num", col = "lightpink")
```

## Boxplot of Occupation-num



```
# Boxplot for continuous variable Y by discrete variable X
boxplot(Y~X, data = data, main = "Boxplot of Age by Occupation",
        col = c("lightblue", "lightpink", "lightyellow", "lightgreen"))
```

## Boxplot of Age by Occupation



*# QUESTION 3*

*# (3abc) Generate the 7 datasets*  
*set.seed(123) # For reproducibility*

*# Simulated datasets*

*poisson\_data <- rpois(500, lambda = 5)*

*# Poisson distribution*

*uniform\_data <- runif(500, min = 3, max = 8)*  
*ibution*

*# Continuous uniform distr*

*exponential\_data <- rexp(500, rate = 7)*

*# Exponential distribution*

*beta\_data <- rbeta(500, shape1 = 2, shape2 = 3)*

*# Beta distribution*

*log\_normal\_data <- rlnorm(500, meanlog = 0, sdlog = 1)*

*# Log-normal distribution*

```
datasets <- list(
  "Discrete X" = X,
  "Continuous Y" = Y,
  "Poisson( $\lambda=5$ )" = poisson_data,
  "Uniform[3,8]" = uniform_data,
  "Exponential( $\lambda=7$ )" = exponential_data,
  "Beta( $\alpha=2$ ,  $\beta=3$ )" = beta_data,
  "Log-Normal( $\mu=0$ ,  $\sigma=1$ )" = log_normal_data
)
```

```
# Draw 100 samples and compute means for n = 10, 50, 100
compute_sample_means <- function(data, n) {
  replicate(100, mean(sample(data, size = n, replace = TRUE)))
}

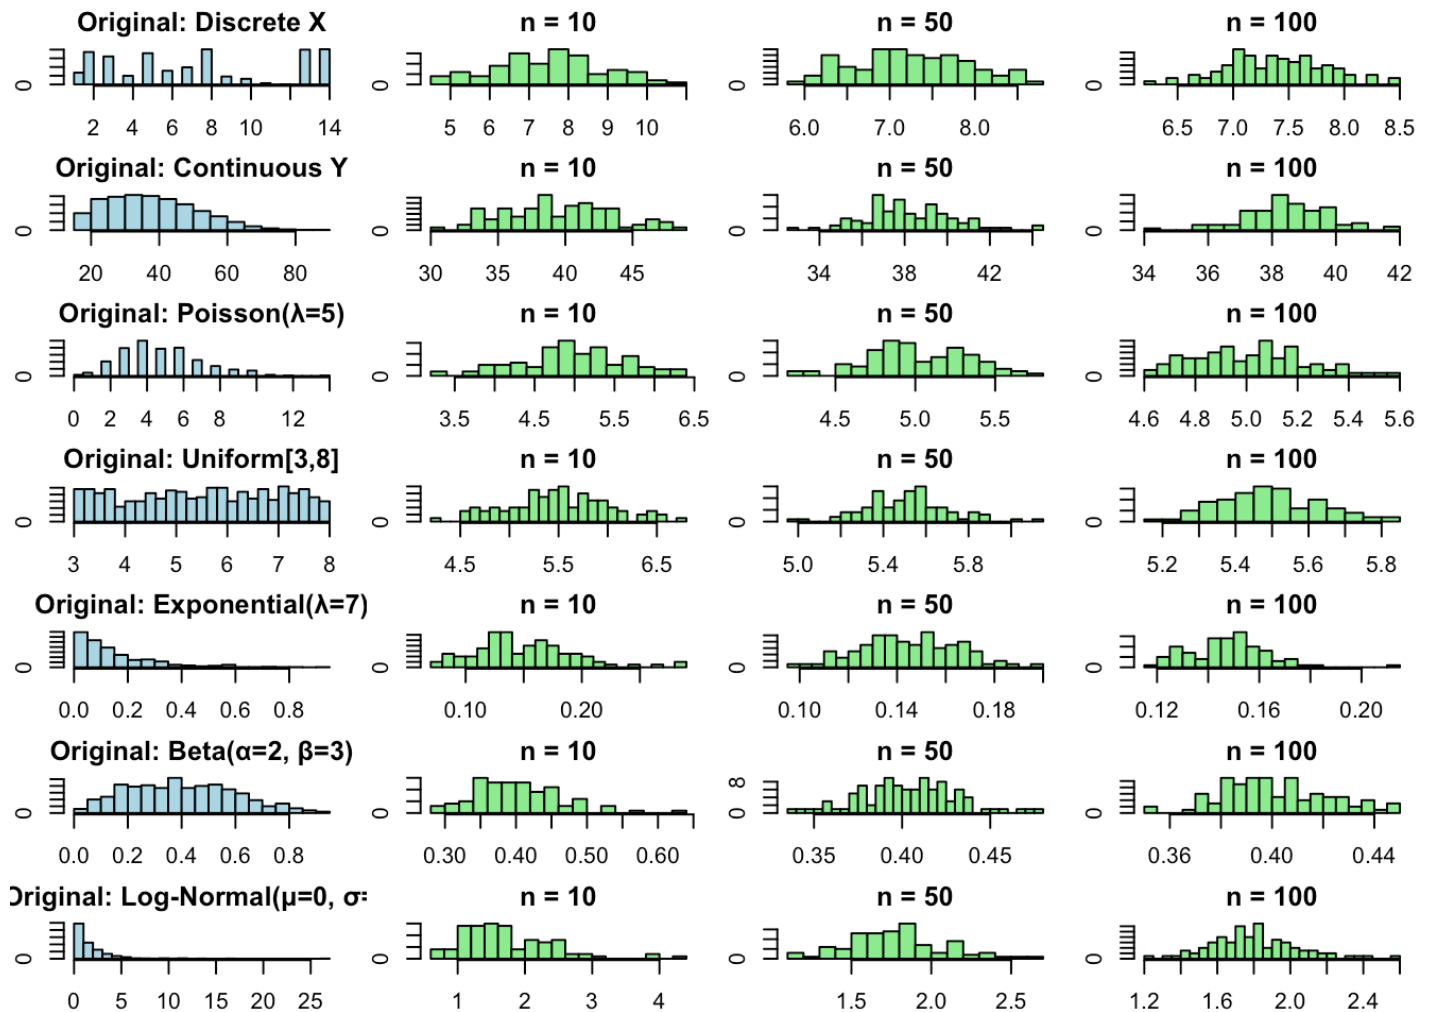
# Compute histograms
plot_histograms <- function(datasets, sample_sizes) {
  par(mfrow = c(7, 4), mar = c(2, 2, 2, 1)) # 7 rows, 4 columns

  for (i in 1:length(datasets)) {
    data <- datasets[[i]]

    # Plot histogram of the original dataset
    hist(data, main = paste("Original:", names(datasets)[i]), col = "lightblue",
         xlab = "Values", ylab = "Frequency", breaks = 20)

    # Compute and plot histograms for sample means
    for (n in sample_sizes) {
      sample_means <- compute_sample_means(data, n)
      hist(sample_means, main = paste("n =", n), col = "lightgreen",
           xlab = "Sample Mean", ylab = "Frequency", breaks = 20)
    }
  }
}

# Run the plotting function for sample sizes n = 10, 50, 100
plot_histograms(datasets, sample_sizes = c(10, 50, 100))
```



# QUESTION 4

# Load required libraries

**library**(missMethods) # For introducing missing data

**library**(mice) # For imputation

##

## Attaching package: 'mice'

## The following object is masked from 'package:stats':

##

## filter

## The following objects are masked from 'package:base':

##

## cbind, rbind

**library**(ggplot2) # For plotting

# Set seed for reproducibility

```

set.seed(123)

# Generate bivariate dataset
n <- 500
X <- rnorm(n, mean = 0, sd = 1) #  $X \sim N(0, 1)$ 
Y <- rnorm(n, mean = 6, sd = 2) #  $Y \sim N(6, 4)$ 
data <- data.frame(X, Y)

# Function to introduce missing data mechanisms
delete_data <- function(data, p, mechanism, ctrl_col = NULL) {
  if (mechanism == "MCAR") {
    return(delete_MCAR(data, p = p, cols_mis = "X"))
  } else if (mechanism == "MAR") {
    return(delete_MAR_rank(data, p = p, cols_mis = "X", cols_ctrl = ctrl_col))
  } else if (mechanism == "NMAR") {
    return(delete_MNAR_censoring(data, p = p, cols_mis = "X"))
  }
}

# Function to plot missing data
plot_missing_data <- function(original, missing, title) {
  # Identify missing indices
  missing_indices <- is.na(missing$X)

  # Create a combined dataset
  data_combined <- data.frame(
    X = original$X,
    Y = original$Y,
    Status = ifelse(missing_indices, "Missing", "Observed")
  )

  # Assign colors
  colors <- c("blue", "red")
  names(colors) <- c("Observed", "Missing")

  # Plot using ggplot2
  ggplot(data_combined, aes(x = X, y = Y, color = Status)) +
    geom_point(alpha = 0.7) +
    scale_color_manual(values = colors) +
    labs(title = title, x = "X", y = "Y") +
    theme_minimal()
}

# Imputation methods
# Mean Imputation
impute_mean <- function(data) {
  data$X[is.na(data$X)] <- mean(data$X, na.rm = TRUE)
  return(data)
}

```



```
# PMM using mice
impute_pmm <- function(data) {
  imputed <- mice(data, method = "pmm", m = 1, maxit = 5, print = FALSE)
  return(complete(imputed))
}

# RMSE calculation
rmse <- function(original, imputed, missing_indices) {
  sqrt(mean((original[missing_indices] - imputed[missing_indices])^2, na.rm = TRUE))
}

# Evaluate imputation methods
evaluate_imputation <- function(data_original, data_missing, method) {
  missing_indices <- is.na(data_missing$X)
  imputed_data <- method(data_missing)
  return(rmse(data_original$X, imputed_data$X, missing_indices))
}

# Initialize results dataframe
results <- data.frame(
  Mechanism = character(),
  Missing_Percentage = numeric(),
  RMSE_Mean = numeric(),
  RMSE_PMM = numeric()
)

# Missingness levels and mechanisms
missingness_levels <- c(0.2, 0.3)
mechanisms <- c("MCAR", "MAR", "NMAR")


# Loop over mechanisms and missingness levels
for (mechanism in mechanisms) {
  for (p in missingness_levels) {
    # Create missing data
    if (mechanism == "MAR") {
      data_missing <- delete_data(data, p, mechanism, ctrl_col = "Y")
    } else {
      data_missing <- delete_data(data, p, mechanism)
    }

    # Imputation and RMSE
    rmse_mean <- evaluate_imputation(data, data_missing, impute_mean)
    rmse_pmm <- evaluate_imputation(data, data_missing, impute_pmm)


    # Append to results
    results <- rbind(
      results,
      data.frame(
        Mechanism = mechanism,
```


```
    Missing_Percentage = p * 100,  
    RMSE_Mean = rmse_mean,  
    RMSE_PMM = rmse_pmm  
  )  
)  
  
# Plot missing data  
title <- paste(mechanism, ": ", p * 100, "% Missing X", sep = "")  
print(plot_missing_data(data, data_missing, title))  
}  
}
```

MCAR: 20% Missing X




MCAR: 30% Missing X







## MAR: 30% Missing X



## NMAR: 20% Missing X



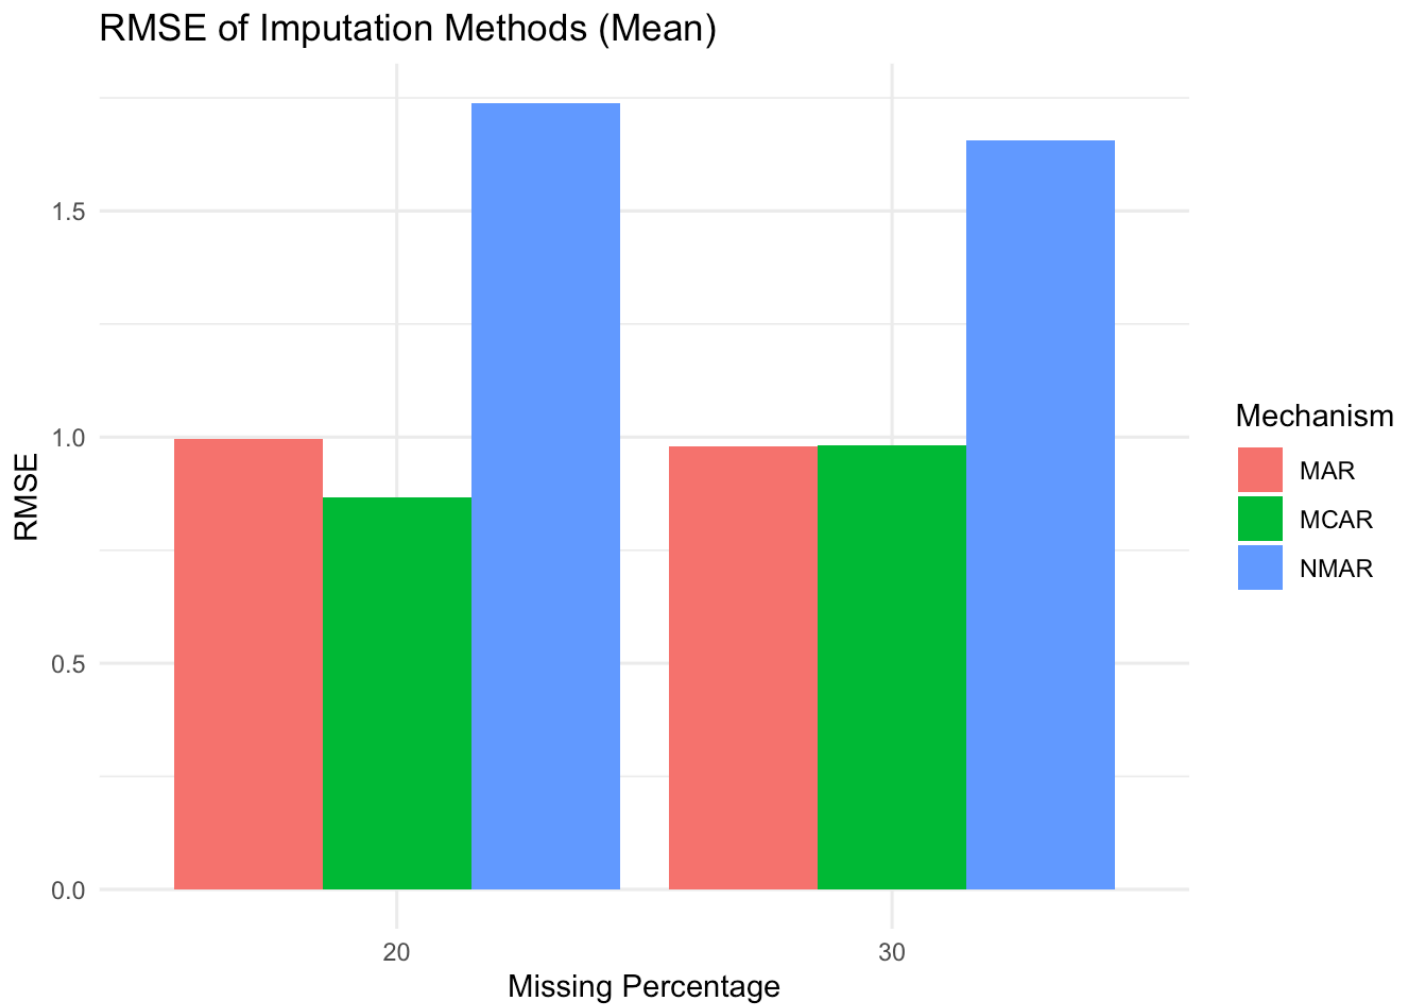
## NMAR: 30% Missing X



```
# Print results
print(results)
```

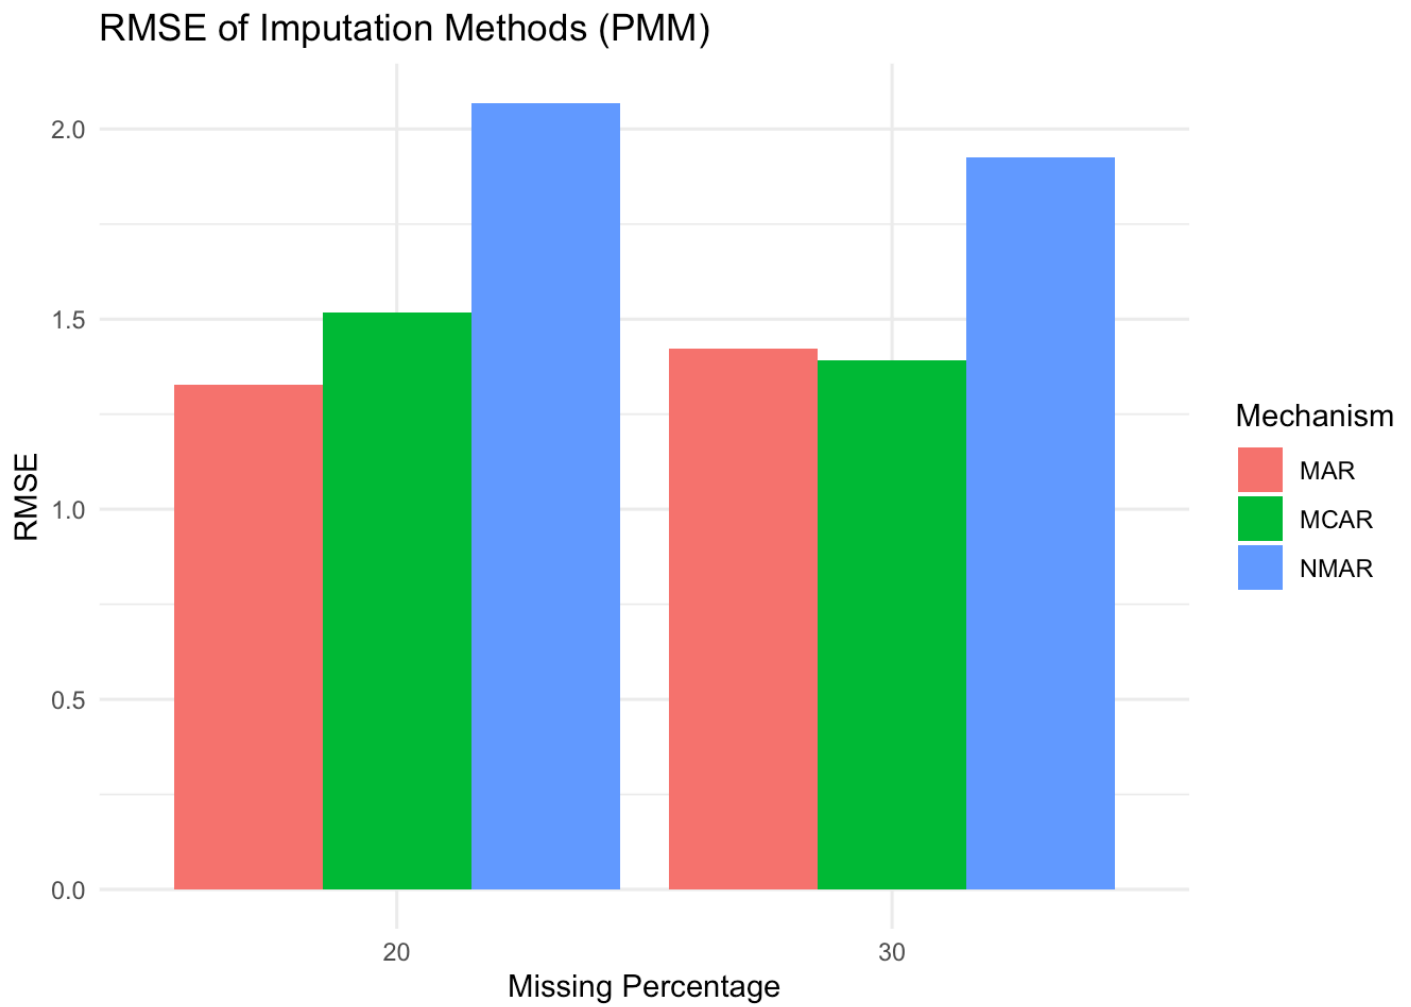
```
##      Mechanism Missing_Percentage RMSE_Mean RMSE_PMM
## 1      MCAR                20 0.8656281 1.516610
## 2      MCAR                30 0.9815377 1.392496
## 3      MAR                 20 0.9963683 1.328518
## 4      MAR                 30 0.9791792 1.422421
## 5      NMAR                20 1.7385654 2.068372
## 6      NMAR                30 1.6548717 1.924810
```

```
# Visualize RMSE results for Mean Imputation
ggplot(results, aes(x = factor(Missing_Percentage), y = RMSE_Mean, fill = Mechanism)) +
  geom_bar(stat = "identity", position = "dodge") +
  labs(title = "RMSE of Imputation Methods (Mean)", x = "Missing Percentage", y = "RMSE") +
  theme_minimal()
```



```
# Visualize RMSE results for PMM
ggplot(results, aes(x = factor(Missing_Percentage), y = RMSE_PMM, fill = Mechanism)) +
  geom_bar(stat = "identity", position = "dodge") +
  labs(title = "RMSE of Imputation Methods (PMM)", x = "Missing Percentage", y = "RMSE") +
  theme_minimal()
```





```
# QUESTION 5
```

```
library(fitdistrplus)
```

```
## Loading required package: MASS
```

```
## Loading required package: survival
```

```
data <- read.csv("/Users/prabuddhadurges/Downloads/adult_data.csv")  
head(data)
```

```
##      age      workclass fnlwgt  education education.num      marital.status
## 1   39      State-gov  77516   Bachelors           13      Never-married
## 2   50  Self-emp-not-inc  83311   Bachelors           13  Married-civ-spouse
## 3   38      Private  215646    HS-grad            9      Divorced
## 4   53      Private  234721     11th             7  Married-civ-spouse
## 5   28      Private  338409   Bachelors           13  Married-civ-spouse
## 6   37      Private  284582   Masters            14  Married-civ-spouse
##      occupation  relationship  race      sex capital.gain capital.loss
## 1   Adm-clerical  Not-in-family  White   Male         2174           0
## 2   Exec-managerial      Husband  White   Male           0           0
## 3  Handlers-cleaners  Not-in-family  White   Male           0           0
## 4  Handlers-cleaners      Husband  Black   Male           0           0
## 5   Prof-specialty      Wife  Black   Female          0           0
## 6   Exec-managerial      Wife  White   Female          0           0
##  hours.per.week native.country income occupation.num
## 1           40  United-States  <=50K           2
## 2           13  United-States  <=50K          13
## 3           40  United-States  <=50K           1
## 4           40  United-States  <=50K           1
## 5           40           Cuba  <=50K          14
## 6           40  United-States  <=50K          13
```

```
A <- data$occupation.num # Discrete variable
B <- data$education.num  # Discrete variable
C <- data$age            # Continuous variable
D <- data$hours.per.week # Continuous variable

# A <- occupation.num (discrete), B <- education.num (discrete), C <- age (continuous), D <- hours.per.week (continuous)

# Step 1: Dataset A - Discrete Variable: occupation.num
A <- data$occupation.num

# (5a) Compute the median of the data.
a <- median(A, na.rm = TRUE) # Compute the median of occupation.num
cat("Median of occupation.num (A):", a, "\n")
```

```
## Median of occupation.num (A): 7
```

```
# (5b) Select a simple random sample of size 100 from the dataset.
set.seed(123) # For reproducibility
sample_A <- sample(A, size = 100, replace = TRUE)

# (5c) Count the number of observations in the sample that are less than 'a'.
Z <- sum(sample_A <= a) # Count how many values are less than or equal to the median
cat("Number of observations in the sample less than or equal to median:", Z, "\n")
```

```
## Number of observations in the sample less than or equal to median: 52
```

```
# (5d) Test the hypothesis that  $p = P[X \leq a]$  is larger than 0.5 at the 5% significance level.
```

```
p0 <- 0.5 # Null hypothesis:  $p = 0.5$ 
```

```
p_hat <- Z / 100 # Proportion of observations less than or equal to median
```

```
# Perform the binomial test
```

```
test_result <- binom.test(Z, 100, p = p0, alternative = "greater", conf.level = 0.95)
```

```
cat("Hypothesis test result for occupation.num:\n")
```

```
## Hypothesis test result for occupation.num:
```

```
print(test_result)
```

```
##
## Exact binomial test
##
## data: Z and 100
## number of successes = 52, number of trials = 100, p-value = 0.3822
## alternative hypothesis: true probability of success is greater than 0.5
## 95 percent confidence interval:
## 0.4332319 1.0000000
## sample estimates:
## probability of success
## 0.52
```

```
# (5e) Provide an approximate 95% confidence interval for p.
```

```
prop_test <- prop.test(Z, 100, conf.level = 0.95)
```

```
cat("95% Confidence Interval for p (occupation.num):\n")
```

```
## 95% Confidence Interval for p (occupation.num):
```

```
print(prop_test$conf.int)
```

```
## [1] 0.4183183 0.6201278
## attr(,"conf.level")
## [1] 0.95
```

```
# QUESTION 6
```

```
B <- data$education.num  
B <- sample(B, size = 3000)
```

```
# (a) Fit a normal distribution to dataset B  
fit_b <- fitdist(B, "norm")  
cat("Fitted Normal Distribution for education.num(B):\n")
```

```
## Fitted Normal Distribution for education.num(B):
```

```
print(fit_b)
```

```
## Fitting of the distribution ' norm ' by maximum likelihood  
## Parameters:  
##      estimate Std. Error  
## mean 10.121333 0.04590066  
## sd    2.514083 0.03245665
```

```
# (b) Conduct a Shapiro-Wilk test for normality  
shapiro_test_b <- shapiro.test(B)  
cat("Shapiro-Wilk Test for Normality education.num(B):\n")
```

```
## Shapiro-Wilk Test for Normality education.num(B):
```

```
print(shapiro_test_b)
```

```
##  
## Shapiro-Wilk normality test  
##  
## data: B  
## W = 0.92561, p-value < 2.2e-16
```

```
set.seed(123)  
data <- data.frame(education.num = rnorm(500, mean = 12, sd = 2)) # Simulated data  
  
# Fit a normal distribution to dataset B (education.num)  
fit <- fitdist(data$education.num, "norm")  
  
# Print the summary of the fitted distribution  
cat("Summary of the fitted normal distribution:\n")
```


```
## Summary of the fitted normal distribution:
```

```
summary(fit)
```


```
## Fitting of the distribution ' norm ' by maximum likelihood
## Parameters :
##      estimate Std. Error
## mean 12.069181 0.08692009
## sd    1.943592 0.06146171
## Loglikelihood: -1041.738   AIC:  2087.476   BIC:  2095.906
## Correlation matrix:
##              mean          sd
## mean  1.000000e+00 -2.277541e-09
## sd    -2.277541e-09  1.000000e+00
```

```
# Visualize the fitted distribution
plot(fit)
```


**Empirical and theoretical dens.**




**Q-Q plot**



**Empirical and theoretical CDFs**



**P-P plot**



```
# Conduct a goodness-of-fit test
gof_results <- gofstat(fit)

# Print the goodness-of-fit test results
cat("Goodness-of-Fit Test Results:\n")
```

```
## Goodness-of-Fit Test Results:
```

```
print(gof_results)
```

```
## Goodness-of-fit statistics
##                                1-mle-norm
## Kolmogorov-Smirnov statistic 0.02300307
## Cramer-von Mises statistic   0.04776675
## Anderson-Darling statistic   0.28007918
##
## Goodness-of-fit criteria
##                                1-mle-norm
## Akaike's Information Criterion 2087.476
## Bayesian Information Criterion 2095.906
```

```
# QUESTION 7
```

```
library(car)
```

```
## Loading required package: carData
```

```
B <- data$education.num

# (a) Divide dataset B into two subsets B1 and B2 (3:2 ratio)
set.seed(123)
sample_indices_B <- sample(1:length(B), size = 0.6 * length(B)) # 60% for B1
B1 <- B[sample_indices_B]
B2 <- B[-sample_indices_B] # Remaining 40% for B2

# (b) Levene's test for equality of variances between B1 and B2

levene_test_B <- leveneTest(c(B1, B2) ~ factor(c(rep(1, length(B1)), rep(2, length(B2)))))
cat("Levene's Test for Equality of Variances between B1 and B2 (hours.per.week):\n")
```

```
## Levene's Test for Equality of Variances between B1 and B2 (hours.per.week):
```

```
print(levene_test_B)
```

```
## Levene's Test for Homogeneity of Variance (center = median)
##           Df F value  Pr(>F)
## group    1  3.3522 0.06771 .
##           498
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
# (c) t-test for equality of means between B1 and B2
t_test_result_B <- t.test(B1, B2)
cat("t-test for Equality of Means between B1 and B2 (hours.per.week):\n")
```

```
## t-test for Equality of Means between B1 and B2 (hours.per.week):
```

```
print(t_test_result_B)
```

```
##
## Welch Two Sample t-test
##
## data:  B1 and B2
## t = -0.95988, df = 395.02, p-value = 0.3377
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
## -0.5307886  0.1825201
## sample estimates:
## mean of x mean of y
## 11.99953 12.17366
```

```
# (d) Confidence Interval for the difference of means (99% CI)
conf_interval_B <- t.test(B1, B2)$conf.int
cat("99% Confidence Interval for the Difference in Means (hours.per.week):\n")
```

```
## 99% Confidence Interval for the Difference in Means (hours.per.week):
```

```
print(conf_interval_B)
```

```
## [1] -0.5307886  0.1825201
## attr(,"conf.level")
## [1] 0.99
```